

A Bayesian Population-Dynamics Model for Clonal T Cells across Tissue and Phenotype in Glioblastoma

June 24, 2026

Contents

1	Setup	2
1.1	Indices	2
1.2	Observed count tensor	2
1.3	Joint tissue–phenotype state space	3
1.4	Valid forward steps	3
1.5	Clone-level observations within a step	3
1.6	Two kinds of destination zero	4
1.7	Setup summary	4
1.8	Notation summary	5
2	Joint Tissue–Phenotype Growth Matrix	5
2.1	Non-negative rows and growth factors	6
2.2	Destination intensity vector	6
2.3	Role of forward steps	6
3	Bayesian Reading of the Growth Matrix	6
3.1	Forward generative reading	6
3.2	Likelihood for destination counts	7
3.3	Prior over growth-matrix entries	7
3.4	Posterior over growth matrices	7
3.5	Posterior predictive destination counts	7
3.6	Backward attribution	8
4	Full Generative Story	8
5	Inference	8
5.1	Learned quantity	10
5.2	Variational approximation	10

5.3	Latent allocation counts	10
5.4	Variational family	10
5.5	Coordinate updates	10
5.5.1	Allocation update	11
5.5.2	Entry update	11
5.6	Optimization objective	11
5.7	Estimated growth matrix	11
5.8	Algorithm summary	11
6	Posterior Summaries	12
6.1	Trafficking and growth	12
6.2	Rate decomposition	12

1 Setup

We observe annotated T cells from multiple patients, anatomical sites, and study timepoints. Each observed cell carries five labels: a patient, a TCR clonotype, a study timepoint, an anatomical site, and a discrete T-cell phenotype state.

The data are represented as aggregate clone-by-timepoint-by-site-by-phenotype counts, from which we construct clone-level observations across forward study steps.

1.1 Indices

We use the following indices throughout.

- Patients: $i = 1, \dots, I$.
- TCR clonotypes within patient i : $c = 1, \dots, C_i$.
- Shared study timepoints: $r \in \{1, \dots, R\}$, ordered as $\tau_1 < \tau_2 < \dots < \tau_R$. The timepoint labels are comparable across patients.
- Anatomical site: $s \in \mathcal{S}$, with $\mathcal{S} = \{\text{PBMC}, \text{CSF}, \text{Tumor}\}$.
- Annotated T-cell phenotype state: $k \in \{1, \dots, K\}$. In the current application $K = 11$, and all phenotypes share a single phenotype space.

1.2 Observed count tensor

For each patient i , clonotype c , timepoint r , site s , and phenotype k , let

$$Y_{icrsk} \in \mathbb{Z}_{\geq 0}$$

be the observed number of cells assigned to that combination: how many cells from patient i , clonotype c , and timepoint r were sampled from site s and assigned phenotype k . Collecting the site and phenotype axes,

$$Y_{icr} \in \mathbb{Z}_{\geq 0}^{|\mathcal{S}| \times K}$$

is the tissue-by-phenotype count tensor for clonotype c in patient i at timepoint r .

Each tissue sample also carries a sequencing depth and a missingness flag. For patient i , timepoint r , and tissue s , let

$$d_{i,r,s} \in \mathbb{Z}_{\geq 0}$$

be the per-sample sequencing depth: the VDJ cell count for the Cell Ranger run of that (patient, timepoint, tissue) sample. It is the Poisson exposure used below. Let

$$m_{i,r,s} \in \{0, 1\}$$

be the missingness indicator, equal to 1 when that sample was profiled and 0 otherwise.

1.3 Joint tissue–phenotype state space

The model operates on the joint tissue–phenotype state space

$$\mathcal{Z} = \mathcal{S} \times \{1, \dots, K\}, \quad z = (s, k),$$

where s is the tissue and k is the phenotype. Since $|\mathcal{S}| = 3$ and $K = 11$, the number of joint states is

$$L = |\mathcal{Z}| = |\mathcal{S}| K = 33.$$

The model therefore tracks clone counts across a 33-state tissue–phenotype space. When describing transitions we write the source state as $z = (a, u)$ and the destination state as $z' = (b, v)$, where $a, b \in \mathcal{S}$ are source and destination tissues and $u, v \in \{1, \dots, K\}$ are source and destination phenotypes.

1.4 Valid forward steps

Forward steps are defined per patient. For patient i , let $\mathcal{R}_i \subseteq \{1, \dots, R - 1\}$ be the set of source timepoints at which patient i has a valid forward step. We index a step by its source timepoint $r \in \mathcal{R}_i$; the step runs from timepoint r to the next sampled timepoint, written $r + 1$. Thus r is not a matrix index—it simply labels a forward step within patient i by where it starts.

Remark (Skipped timepoints). Writing the destination as $r + 1$ assumes consecutive sampling. When a patient is missing a timepoint—so that, say, T_1 and T_3 are paired— $r + 1$ should be read as the *next available* timepoint for that patient. Which timepoints are paired is a data-preprocessing detail and does not change anything below.

Each $r \in \mathcal{R}_i$ is one observed forward step. Steps such as $T_1 \rightarrow T_2$, $T_2 \rightarrow T_3$, and $T_3 \rightarrow T_4$ are treated as repeated observations of the *same* forward process: the growth matrix is shared across all patients and steps, and the indices i and r merely organize observations rather than selecting a separate matrix.

1.5 Clone-level observations within a step

Fix a patient i , a valid step $r \in \mathcal{R}_i$, and a clonotype c observed for that patient. The source and destination count vectors over the joint state space are read directly off the count tensor at timepoints r and $r + 1$,

$$x_{irc}(s, k) = Y_{i c r s k}, \quad y_{irc}(s, k) = Y_{i c (r+1) s k}, \quad (s, k) \in \mathcal{Z},$$

so $x_{irc} \in \mathbb{Z}_{\geq 0}^L$ and $y_{irc} \in \mathbb{Z}_{\geq 0}^L$.

The source enters the model only after dividing out the source sequencing depth. The *depth-rescaled source* is

$$\tilde{x}_{irc}(z) = \frac{x_{irc}(z)}{d_{ir,s}^{\text{src}}}, \quad s = \text{tissue}(z), \quad d_{ir,s}^{\text{src}} = d_{i,r,s},$$

i.e. each source count is divided by the source-timepoint depth of its tissue. It is treated as a fixed, known plug-in, not a random quantity: source abundance enters after dividing out source depth, and residual source-sampling variation is not modeled.

We retain clonotypes present at the *source* end of the step,

$$C_{ir} = \{c \in \{1, \dots, C_i\} : N_{irc}^{\text{src}} > 0\}, \quad N_{irc}^{\text{src}} = \sum_{z \in \mathcal{Z}} x_{irc}(z).$$

A destination total of zero is now a modeled observation—a contraction or death signal—not an exclusion criterion. This gives the data hierarchy

$$i \longrightarrow r \in \mathcal{R}_i \longrightarrow c \in C_{ir},$$

patients contain valid forward steps, and each forward step has a retained set of source-present clonotypes. A *clone-level forward observation* is the triple (i, r, c) ; for compact notation we write

$$j \equiv (i, r, c), \quad x_j = x_{irc}, \quad \tilde{x}_j = \tilde{x}_{irc}, \quad y_j = y_{irc}.$$

The destination vector y_j keeps its per-tissue structure: each tissue sub-block $\{(s', k) : k = 1, \dots, K\}$ of y_j carries its own destination depth $d_{j,s'} = d_{i,r+1,s'}$ and missingness flag $m_{j,s'} = m_{i,r+1,s'}$.

1.6 Two kinds of destination zero

Destination zeros are part of the likelihood specification, and the two kinds are treated differently.

1. **Missing sample** ($m_{j,s'} = 0$). The tissue sub-block $\{(s', k)\}$ was not profiled at the destination timepoint. It is excluded from the likelihood entirely—masked out, never entered as a zero count.
2. **Depth-zero** ($m_{j,s'} = 1$, clone not captured). A genuine observed zero, modeled natively by the Poisson likelihood below: a small mean $d_{j,s'} \mu_j(z')$ emits zero with positive probability while $\mu_j(z') > 0$. No zero-inflation component is added.

The depth-zero mechanism is what permits a clone to be absent at one step and present at the next.

1.7 Setup summary

The setup produces the clone-level forward observations

$$\mathcal{D} = \{(\tilde{x}_j, y_j, d_j, m_j, i_j, r_j, c_j)\}_{j=1}^J, \quad j \equiv (i_j, r_j, c_j),$$

each carrying a depth-rescaled source vector, a destination count vector, the per-tissue destination depths $d_j = (d_{j,s'})_{s' \in \mathcal{S}}$ and missingness flags $m_j = (m_{j,s'})_{s' \in \mathcal{S}}$, and a patient / step / clonotype identity. The count vectors live in the joint tissue–phenotype state space $\mathcal{Z} = \mathcal{S} \times \{1, \dots, K\}$, with $|\mathcal{Z}| = 33$.

1.8 Notation summary

Table 1 collects the core modeling symbols; the variational quantities used only for fitting (the Gamma parameters $\text{shape}_{zz'}$, $\text{rate}_{zz'}$, the responsibilities $r_{jzz'}$, and ψ) are introduced where they are used in the inference section.

Symbol	Meaning
$i = 1, \dots, I$	patient index
$c = 1, \dots, C_i$	TCR clonotype within patient i
$r \in \{1, \dots, R\}$	shared study timepoint, $\tau_1 < \dots < \tau_R$
$s \in \mathcal{S}$	anatomical site, $\mathcal{S} = \{\text{PBMC}, \text{CSF}, \text{Tumor}\}$
$k \in \{1, \dots, K\}$	phenotype state ($K = 11$)
$z = (s, k) \in \mathcal{Z}$	joint tissue–phenotype state ($L = \mathcal{Z} = 33$)
Y_{icrsk}	observed count tensor
$d_{i,r,s}$	per-sample sequencing depth (Poisson exposure)
$m_{i,r,s} \in \{0, 1\}$	sample missingness indicator
$\mathcal{R}_i$	valid forward steps for patient i ($\subseteq \{1, \dots, R-1\}$)
$r \in \mathcal{R}_i$	forward step, from source timepoint r to $r+1$
C_{ir}	clonotypes present at the source end of step r
$j = (i_j, r_j, c_j)$	clone-level observation, $j = 1, \dots, J$
$x_j, y_j \in \mathbb{Z}_{\geq 0}^L$	source / destination count vectors
$\tilde{x}_j$	depth-rescaled source, $\tilde{x}_j(z) = x_j(z) / d_{j,s}^{\text{src}}$
$M \in \mathbb{R}_{\geq 0}^{L \times L}$	non-negative growth matrix (free row sums)
$\mu_j = \tilde{x}_j M \in \mathbb{R}_{\geq 0}^L$	destination intensity vector
a_z, b_z	Gamma prior shape / rate for row z
$w_{jzz'}$	backward attribution probability
$\xi_{jzz'}$	latent allocation count (computational)
$\mathcal{D}$	full dataset

Table 1: Core modeling notation.

2 Joint Tissue–Phenotype Growth Matrix

The central latent object is a single joint tissue–phenotype *growth matrix*

$$M \in \mathbb{R}_{\geq 0}^{L \times L}, \quad L = 33,$$

whose rows are source states and whose columns are destination states. For source state $z = (a, u)$ and destination state $z' = (b, v)$, the entry $M_{zz'}$ is the expected number of destination- z' cells produced per source- z cell over one forward step (one 8-week interval):

$$M_{zz'} = M_{(a,u),(b,v)} = \mathbb{E}[\text{cells in state } (b, v) \text{ at } r + 1 \text{ per cell in state } (a, u) \text{ at } r].$$

2.1 Non-negative rows and growth factors

M has non-negative entries, $M_{zz'} \geq 0$, and *free* row sums—rows are not constrained to sum to one. The row sum

$$g_z = \sum_{z' \in \mathcal{Z}} M_{zz'}$$

is the net 8-week growth factor of source state z : $g_z > 1$ is net expansion of that state, $g_z < 1$ net contraction or death.

2.2 Destination intensity vector

For observation j with depth-rescaled source $\tilde{x}_j$, the destination *intensity vector* is the push-forward of $\tilde{x}_j$ through M :

Central growth calculation

$$\mu_j = \tilde{x}_j M \in \mathbb{R}_{\geq 0}^L, \quad \text{equivalently} \quad \mu_j(z') = \sum_{z \in \mathcal{Z}} \tilde{x}_j(z) M_{zz'}.$$

The expected destination count in state z' scales the intensity by the destination depth of its tissue,

$$\mathbb{E}[y_j(z') \mid M] = d_{j,s'} \mu_j(z'), \quad s' = \text{tissue}(z').$$

There is no simplex constraint on μ_j and no conditioning on a destination total.

2.3 Role of forward steps

The same matrix M is used for every patient i and step $r \in \mathcal{R}_i$. The pair (i, r) selects which source and destination vectors to build, not a different matrix; all steps are repeated observations of a common one-step forward process.

3 Bayesian Reading of the Growth Matrix

The matrix M is a latent non-negative kernel of per-cell destination intensities over the joint tissue-phenotype state space.

3.1 Forward generative reading

Read forward, the model is a Poisson superposition. Treating the depth-rescaled source as a population with $\tilde{x}_j(z)$ cells in each state z , each source cell in state z independently emits a Poisson-distributed number of cells into each destination state z' with mean $M_{zz'}$. Summing

these emissions over all source states and scaling by the destination depth, the destination- z' count is, by Poisson superposition,

$$y_j(z') \mid M \sim \text{Poisson}\left(d_{j,s'} \sum_{z \in \mathcal{Z}} \tilde{x}_j(z) M_{zz'}\right) = \text{Poisson}(d_{j,s'} \mu_j(z')).$$

3.2 Likelihood for destination counts

Observation likelihood

$$y_j(z') \sim \text{Poisson}(d_{j,s'} (\tilde{x}_j M)(z')), \quad s' = \text{tissue}(z'),$$

evaluated only over non-missing destination sub-blocks ($m_{j,s'} = 1$). The latent matrix enters the likelihood only through the intensity $\mu_j = \tilde{x}_j M$.

3.3 Prior over growth-matrix entries

Each entry receives an independent Gamma prior in shape–rate form,

$$M_{zz'} \sim \text{Gamma}(a_z, b_z),$$

with row-specific shape a_z and rate b_z .

Remark (Support restrictions are an implementation detail). In practice the implementation may exclude biologically implausible transitions by fixing the corresponding entries of M to zero (placing the Gamma prior only on the retained entries of each row). This sparsifies M but does not change the form of the model; sums over destination states below are read over the retained entries. We treat the support as a fixed preprocessing choice.

For the likelihood to support the data, the support guard is: for every non-missing destination state z' with $y_j(z') > 0$, at least one source state in the support of $\tilde{x}_j$ must reach z' (i.e. have $M_{zz'} > 0$).

3.4 Posterior over growth matrices

After observing the data $\mathcal{D}$, inference targets the posterior $p(M \mid \mathcal{D})$, which captures uncertainty about the joint tissue–phenotype growth kernel after all clone-level forward observations. Posterior summaries include posterior means $\mathbb{E}[M_{zz'} \mid \mathcal{D}]$ and credible intervals for individual entries.

3.5 Posterior predictive destination counts

Because M is latent, the intensity $\mu_j = \tilde{x}_j M$ is also uncertain. Posterior draws of M induce replicated destination counts

$$y_j^{\text{rep}}(z') \sim \text{Poisson}(d_{j,s'} \mu_j(z'))$$

over the non-missing sub-blocks ($m_{j,s'} = 1$).

3.6 Backward attribution

The reverse question—which source state produced an observed destination cell—is a posterior attribution. For observation j ,

$$w_{jzz'} = \frac{\tilde{x}_j(z) M_{zz'}}{\sum_{\tilde{z} \in \mathcal{Z}} \tilde{x}_j(\tilde{z}) M_{\tilde{z}z'}} = \frac{\tilde{x}_j(z) M_{zz'}}{\mu_j(z')},$$

the same form as the forward intensity with source and destination roles reversed. Because the destination counts are Poisson, these probabilities also induce a multinomial decomposition of each observed count $y_j(z')$ into latent source-to-destination allocation counts. Those counts are used only for fitting, so we introduce them together with the inference algorithm in Section 5.

4 Full Generative Story

The full model is a model for the one-step expansion and redistribution of clone mass across tissue–phenotype states. It can be summarized as follows.

Generative model

1. **State space.** Define $\mathcal{Z} = \mathcal{S} \times \{1, \dots, K\}$ with $|\mathcal{Z}| = 33$. Each state $z = (s, k)$ pairs an anatomical tissue with an annotated T-cell phenotype.
2. **Growth-matrix entries.** For each source state z and destination state z' , draw independently
$$M_{zz'} \sim \text{Gamma}(a_z, b_z).$$
Entry $M_{zz'}$ is the expected number of destination- z' cells per source- z cell over one forward step.
3. **Source.** For each patient i , step $r \in \mathcal{R}_i$, and retained clonotype $c \in C_{ir}$, form the observation $j = (i, r, c)$ and the fixed depth-rescaled source $\tilde{x}_j$.
4. **Intensity.** Push the source forward, $\mu_j = \tilde{x}_j M$.
5. **Observation.** For each non-missing destination sub-block ($m_{j,s'} = 1$), draw the destination counts

$$y_j(z') \sim \text{Poisson}(d_{j,s'} \mu_j(z')), \quad s' = \text{tissue}(z').$$


Putting the pieces together, the generative model is, compactly,

$$M_{zz'} \sim \text{Gamma}(a_z, b_z), \quad \mu_j = \tilde{x}_j M, \quad y_j(z') \sim \text{Poisson}(d_{j,s'} \mu_j(z')) \quad (m_{j,s'} = 1),$$

with the corresponding dependency structure shown in Figure 1.

5 Inference

Inference estimates the posterior over the shared growth matrix M . The observed data are the clone-level forward observations $\mathcal{D}$; the depth-rescaled sources $\tilde{x}_j$, the depths $d_{j,s'}$, and the missingness flags $m_{j,s'}$ are fixed inputs, as are the Gamma prior parameters a_z, b_z .



Node	Meaning
x_{irc}	observed source counts at the start of the step (over the 33 states)
$\tilde{x}_{irc} = x_{irc} / d_{irs}^{\text{src}}$	depth-rescaled source (deterministic; folded into the $x \rightarrow \mu$ edge)
$\mu_{irc} = \tilde{x}_{irc} M$	destination intensity vector
y_{irc}	observed destination counts in tissue s at the end of the step
M_z	latent growth-matrix row, shared across all patients and steps
d_{irs}	destination sequencing depth, Poisson exposure (fixed)
m_{irs}	sample missingness flag (fixed)
(a_z, b_z)	Gamma prior shape / rate (fixed)

Figure 1: Plate diagram of the generative model (node key above). The rows M_z are drawn once per source state z from a Gamma prior (a_z, b_z) ; this block sits *outside* the patient plate, so a single M is shared across all patients and steps. The observation side nests patient i , step $r \in \mathcal{R}_i$, clone $c \in \mathcal{C}_{ir}$, and tissue $s \in \mathcal{S}$. The depth-rescaled source $\tilde{x}_{irc}$ is pushed through M to the intensity $\mu_{irc} = \tilde{x}_{irc} M$, and within each tissue sub-block the destination counts y_{irc} are a Poisson draw with mean $d_{irs} \mu_{irc}$, gated by missingness m_{irs} . Double circles (x_{irc} , y_{irc}) mark observed quantities; solid dots ((a_z, b_z) , d_{irs} , m_{irs}) mark fixed inputs; the depth-rescaling $\tilde{x}_{irc} = \text{Rescale}(x_{irc})$ is folded into the $x \rightarrow \mu$ edge. Latent allocation counts used for fitting are a computational augmentation introduced in Section 5.

5.1 Learned quantity

The main learned object is the posterior $p(M \mid \mathcal{D})$. A single-entry summary is

$$\mathbb{E}[M_{(a,u),(b,v)} \mid \mathcal{D}],$$

the posterior mean number of destination- (b, v) cells produced per source- (a, u) cell over one forward step.

5.2 Variational approximation

We approximate $p(M \mid \mathcal{D})$ by mean-field variational inference. The difficulty is that an observed destination count in state z' is a *superposition* of contributions from every source state z , and these source-to-destination allocations are unobserved. Variational inference augments the model with latent allocation counts and fits a factorized approximation by coordinate ascent.

5.3 Latent allocation counts

Let $\xi_{jzz'}$ be the unobserved number of destination- z' cells in observation j contributed by source state z . By Poisson superposition each source contribution is itself Poisson, $\xi_{jzz'} \sim \text{Poisson}(d_{j,s'} \tilde{x}_j(z) M_{zz'})$ independently, and they decompose each observed destination count,

$$\sum_{z \in \mathcal{Z}} \xi_{jzz'} = y_j(z').$$

Conditional on the total $y_j(z')$, the allocations are multinomial with probabilities equal to the backward attribution $w_{jzz'}$ of Section 3.6; variational inference replaces these with an approximate responsibility $r_{jzz'}$.

5.4 Variational family

We adopt a mean-field factorization $q(M, \xi) = q(M) q(\xi)$. The growth matrix receives independent Gamma factors, one per entry,

$$q(M) = \prod_{z \in \mathcal{Z}} \prod_{z' \in \mathcal{Z}} \text{Gamma}(M_{zz'}; \text{shape}_{zz'}, \text{rate}_{zz'}),$$

and the allocation variables receive multinomial factors (a Poisson vector conditioned on its total),

$$q(\xi) = \prod_{j=1}^J \prod_{z' : m_{j,s'}=1} \text{Multinomial}(\xi_{j \cdot z'}; y_j(z'), r_{j \cdot z'}),$$

where $r_{jzz'}$ is the variational responsibility assigning destination state z' in observation j to source state z .

5.5 Coordinate updates

The algorithm alternates between updating the allocation responsibilities $r_{jzz'}$ and the Gamma parameters ($\text{shape}_{zz'}, \text{rate}_{zz'}$).

5.5.1 Allocation update

For each observation j , destination state z' with $m_{j,s'} = 1$, and source state z ,

$$r_{jzz'} \propto \tilde{x}_j(z) \exp(\mathbb{E}_q[\log M_{zz'}]),$$

normalized over source states so that $\sum_{z \in \mathcal{Z}} r_{jzz'} = 1$ for each j and z' (the common depth $d_{j,s'}$ cancels in the normalization). Under the Gamma factor, the required expectation is

$$\mathbb{E}_q[\log M_{zz'}] = \psi(\text{shape}_{zz'}) - \log(\text{rate}_{zz'}),$$

where ψ is the digamma function.

5.5.2 Entry update

For each source state z and destination state z' , the Gamma–Poisson conjugate update is

$$\text{shape}_{zz'} = a_z + \sum_j y_j(z') r_{jzz'}, \quad \text{rate}_{zz'} = b_z + \sum_j d_{j,s'} \tilde{x}_j(z),$$

where both sums run over observations whose destination sub-block s' is non-missing ($m_{j,s'} = 1$). The shape accumulates the expected number of $z \rightarrow z'$ cells allocated from the data; the rate accumulates the depth-weighted source exposure.

5.6 Optimization objective

The coordinate updates maximize the evidence lower bound (ELBO)

$$\mathcal{L}(q) = \mathbb{E}_q[\log p(M, \xi, \mathcal{D})] - \mathbb{E}_q[\log q(M, \xi)],$$

and the algorithm iterates between the allocation update and the entry update until $\mathcal{L}(q)$ stabilizes.

5.7 Estimated growth matrix

After convergence, the fitted growth matrix is the variational posterior mean of each Gamma factor,

$$\hat{M}_{zz'} = \mathbb{E}_q[M_{zz'}] = \frac{\text{shape}_{zz'}}{\text{rate}_{zz'}},$$

explicitly *not* normalized across destinations z' . The matrix $\hat{M}$ is the estimated one-step growth matrix over the 33-state tissue–phenotype space.

5.8 Algorithm summary

Variational algorithm (CAVI)

1. **Initialize** the variational Gamma parameters (shape, rate).
2. **Repeat until the ELBO $\mathcal{L}(q)$ converges:**

- update the allocation responsibilities $r_{jzz'} \propto \tilde{x}_j(z) \exp(\mathbb{E}_q[\log M_{zz'}])$;
- update the entry parameters $\text{shape}_{zz'} = a_z + \sum_j y_j(z') r_{jzz'}$, $\text{rate}_{zz'} = b_z + \sum_j d_{j,s'} \tilde{x}_j(z)$.

3. **Return** the posterior mean $\hat{M} = \mathbb{E}_q[M]$.

6 Posterior Summaries

With the variational posterior fitted and $\hat{M}$ in hand, the following quantities are read off it.

6.1 Trafficking and growth

The fitted matrix $\hat{M}$ admits two readings, on its off-diagonal tissue blocks and on its row sums.

Trafficking. For source state $z = (a, u)$ and a destination tissue $b \neq a$, the off-diagonal block sum

$$\tau_{z \rightarrow b} = \sum_{v=1}^K \hat{M}_{(a,u),(b,v)}$$

is the expected number of progeny of a source- z cell that land in tissue b over the step; summing over destination tissues $b \neq a$ gives the total out-of-tissue trafficking from state z .

Growth. The full row sum (diagonal block included)

$$g_z = \sum_{z' \in \mathcal{Z}} \hat{M}_{zz'}$$

is the net per-cell progeny count of source state z over the step.

Off-diagonal block sums measure trafficking; the row sum, with the diagonal included, measures net growth.

6.2 Rate decomposition

To express the step as per-week rates, define the generator

$$B = \frac{1}{\Delta t} \log \hat{M}, \quad \Delta t = 8 \text{ weeks},$$

the matrix logarithm of $\hat{M}$ divided by the step length. The off-diagonal entries $B_{zz'} \geq 0$ ($z' \neq z$) are the per-week trafficking rates between states, and the row sums of B are the net per-capita growth rates per week. As above, the off-diagonal entries carry trafficking and the row sums carry growth.